

CAS 741 (Development of Scientific Computing Software)

Winter 2025

MPIR Implementations

Xunzhou (Joe) Ye

Faculty of Engineering, McMaster University

April 1, 2025



General Information

- MPIR is a sparse linear solver designed to solve large, sparse real matrices efficiently.
- It uses the General Minimal Residual (GMRES) method for internal matrix solves and iterative refinement techniques to improve both speed and accuracy.
- Intended for use in computational science, engineering, and numerical analysis applications.
- As a complete library suite, the software also includes example programs to demonstrate the solver interfaces and practical use cases of the solver.

Inputs

| Variable | Description |
|-------------------|---|
| A | $n \times n$ matrix |
| b | n -vector |
| ϵ | a solution is found if the norm of the residual is less than ϵ |
| n_{iter} | the maximum number of iterations to perform |
| u_f | factorization precision |
| u_w | working precision |
| u_r | precision in which the residuals are computed |

Assumptions and Constraints

- A1 Matrix \mathbf{A} is symmetric quasi-definite.
- A2 \mathbf{A} is stored in *Compressed Sparse Column Format (CSC)* — *Scipy lecture notes 2025* format.
- A3 Only the upper triangular part of \mathbf{A} is stored.
- A4 The precisions follow the order $u_f \leq u_w \leq u_r$, with u_r being the highest precision.
- C1 Use preconditioned GMRES for solving the error correction vector.

Core Algorithms

Algorithm Iterative refinement

- 1: **for** $m \leftarrow 1, 2, \dots$, the m th iteration **do**
 - 2: $\mathbf{r}_m \leftarrow \mathbf{b} - \mathbf{A}\mathbf{x}_m$ ▷ Compute the residuals
 - 3: Solve $\mathbf{A}\mathbf{d}_m = \mathbf{r}_m$ for \mathbf{d}_m ▷ Compute the correction
 - 4: $\mathbf{x}_{m+1} = \mathbf{x}_m + \mathbf{d}_m$ ▷ Add the correction
 - 5: **end for**
-

Core Algorithms

Algorithm GMRES-IR with \mathbf{LDL}^\top factorization in MP

- 1: Perform \mathbf{LDL}^\top factorization of \mathbf{A} ▷ at u_f
 - 2: Solve $\mathbf{LDL}^\top \mathbf{x}_0 = \mathbf{b}$ ▷ at u_f
 - 3: **for** $i \leftarrow 0, 1, \dots, n_{\text{iter}}$ and $\|r_i\|_2 \geq \epsilon$ **do**
 - 4: $r_i \leftarrow \mathbf{b} - \mathbf{A}\mathbf{x}_i$ ▷ at u_r
 - 5: GMRES Solve $(\mathbf{LDL}^\top)^{-1}\mathbf{A}\mathbf{d}_i = (\mathbf{LDL}^\top)^{-1}\mathbf{r}_i$ ▷ at u_w
 - 6: $\mathbf{x}_{i+1} = \mathbf{x}_i + \mathbf{d}_i$ ▷ at u_r
 - 7: **end for**
-

Floating Point Precisions

| Arithmetic | Sym. | Bits | | Type |
|------------|------|------|------|------------------------|
| | | Sig. | Exp. | |
| bfloat16 | B | 8 | 8 | std::bfloat16_t |
| fp16 | H | 11 | 5 | std::float16_t |
| fp32 | S | 24 | 8 | std::float32_t, float |
| fp64 | D | 53 | 11 | std::float64_t, double |
| fp128 | Q | 113 | 15 | std::float128_t |

Demo

- C++ templates: a bit of metaprogramming to ensure $u_f \leq u_w \leq u_r$.
- CMake/CTest: portability, project dependencies, unit test driver, CI/CD.
- clang-format: also with CI/CD.

Future Work

1. Performance testing
2. Hardly converges in low precisions

References



Compressed Sparse Column Format (CSC) — Scipy lecture notes (2025). URL: https://scipy-lectures.org/advanced/scipy_sparse/csc_matrix.html (visited on 03/26/2025).

Questions